

## ASSIGNMENT-3

### 1:Wordmedian:

```
[acadgild@localhost ~]$ hadoop jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmedian /new.txt wordmedian
19/01/12 21:05:35 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
19/01/12 21:05:38 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
19/01/12 21:05:41 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
19/01/12 21:05:43 INFO input.FileInputFormat: Total input paths to process : 1
19/01/12 21:05:43 INFO mapreduce.JobSubmitter: number of splits:1
19/01/12 21:05:44 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1547306488017_0004
19/01/12 21:05:45 INFO impl.YarnClientImpl: Submitted application application_1547306488017_0004
19/01/12 21:05:46 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1547306488017_0004/
19/01/12 21:05:46 INFO mapreduce.Job: Running job: job_1547306488017_0004
19/01/12 21:06:09 INFO mapreduce.Job: Job job_1547306488017_0004 running in uber mode : false
19/01/12 21:06:09 INFO mapreduce.Job: map 0% reduce 0%
19/01/12 21:06:26 INFO mapreduce.Job: map 100% reduce 0%
19/01/12 21:06:46 INFO mapreduce.Job: map 100% reduce 100%
19/01/12 21:06:48 INFO mapreduce.Job: Job job_1547306488017_0004 completed successfully
19/01/12 21:06:49 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=106
    FILE: Number of bytes written=215283
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=248
    HDFS: Number of bytes written=43
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
```

Hbase

```
[acadgild@localhost ~]$ hadoop jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmedian /new.txt wordmedian
19/01/12 21:05:35 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
19/01/12 21:05:38 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
19/01/12 21:05:41 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
19/01/12 21:05:43 INFO input.FileInputFormat: Total input paths to process : 1
19/01/12 21:05:43 INFO mapreduce.JobSubmitter: number of splits:1
19/01/12 21:05:44 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1547306488017_0004
19/01/12 21:05:45 INFO impl.YarnClientImpl: Submitted application application_1547306488017_0004
19/01/12 21:05:46 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1547306488017_0004/
19/01/12 21:05:46 INFO mapreduce.Job: Running job: job_1547306488017_0004
19/01/12 21:06:09 INFO mapreduce.Job: Job job_1547306488017_0004 running in uber mode : false
19/01/12 21:06:09 INFO mapreduce.Job: map 0% reduce 0%
19/01/12 21:06:26 INFO mapreduce.Job: map 100% reduce 0%
19/01/12 21:06:46 INFO mapreduce.Job: map 100% reduce 100%
19/01/12 21:06:48 INFO mapreduce.Job: Job job_1547306488017_0004 completed successfully
19/01/12 21:06:49 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=106
    FILE: Number of bytes written=215283
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=248
    HDFS: Number of bytes written=43
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
```

Hbase

## 2.WORDMEAN:

```
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost sbin]$ hadoop jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmean /new.txt wordmean
19/01/12 21:10:19 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
19/01/12 21:10:22 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
19/01/12 21:10:26 INFO input.FileInputFormat: Total input paths to process : 1
19/01/12 21:10:27 INFO mapreduce.JobSubmitter: number of splits:1
19/01/12 21:10:27 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1547306488017_0005
19/01/12 21:10:28 INFO impl.YarnClientImpl: Submitted application application_1547306488017_0005
19/01/12 21:10:29 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1547306488017_0005/
19/01/12 21:10:29 INFO mapreduce.Job: Running job: job_1547306488017_0005
19/01/12 21:10:52 INFO mapreduce.Job: Job job_1547306488017_0005 running in uber mode : false
19/01/12 21:10:52 INFO mapreduce.Job: map 0% reduce 0%
19/01/12 21:11:11 INFO mapreduce.Job: map 100% reduce 0%
19/01/12 21:11:29 INFO mapreduce.Job: map 100% reduce 100%
19/01/12 21:11:30 INFO mapreduce.Job: Job job_1547306488017_0005 completed successfully
19/01/12 21:11:31 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=39
    FILE: Number of bytes written=215409
    FILE: Number of read operations=0
    FILE: Number of write operations=0
```

```
Combine Input Records=32
Combine output records=2
Reduce input groups=2
Reduce shuffle bytes=39
Reduce input records=2
Reduce output records=2
Spilled Records=4
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=405
CPU time spent (ms)=3940
Physical memory (bytes) snapshot=293048320
Virtual memory (bytes) snapshot=4118220800
Total committed heap usage (bytes)=170004480

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=154
File Output Format Counters
  Bytes Written=20

The mean is: 4.923076923076923
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost sbin]$
```

### 3:WordStandardDeviation:

```
[acadgild@localhost sbin]$ hadoop jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordstandarddeviation /new.txt wordstandarddeviation
19/01/12 21:14:17 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
19/01/12 21:14:21 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
19/01/12 21:14:25 INFO input.FileInputFormat: Total input paths to process : 1
19/01/12 21:14:26 INFO mapreduce.JobSubmitter: number of splits:1
19/01/12 21:14:26 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1547306488017_0006
19/01/12 21:14:27 INFO impl.YarnClientImpl: Submitted application application_1547306488017_0006
19/01/12 21:14:28 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1547306488017_0006/
19/01/12 21:14:28 INFO mapreduce.Job: Running job: job_1547306488017_0006
19/01/12 21:14:52 INFO mapreduce.Job: Job job_1547306488017_0006 running in uber mode : false
19/01/12 21:14:52 INFO mapreduce.Job: map 0% reduce 0%
19/01/12 21:15:11 INFO mapreduce.Job: map 100% reduce 0%
19/01/12 21:15:30 INFO mapreduce.Job: map 100% reduce 100%
19/01/12 21:15:31 INFO mapreduce.Job: Job job_1547306488017_0006 completed successfully
19/01/12 21:15:32 INFO mapreduce.Job: Counters: 49
```

#### File System Counters

```
FILE: Number of bytes read=56
FILE: Number of bytes written=215629
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=248
HDFS: Number of bytes written=31
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
```

#### Job Counters

#### Map-Reduce Framework

```
Map input records=1
Map output records=78
Map output bytes=1144
Map output materialized bytes=56
Input split bytes=94
Combine input records=78
Combine output records=3
Reduce input groups=3
Reduce shuffle bytes=56
Reduce input records=3
Reduce output records=3
Spilled Records=6
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=410
CPU time spent (ms)=3930
Physical memory (bytes) snapshot=293691392
Virtual memory (bytes) snapshot=4123090944
Total committed heap usage (bytes)=170004480
```

#### Shuffle Errors

```
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
```

#### File Input Format Counters

```
Bytes Read=154
```

#### File Output Format Counters

```
Bytes Written=31
```

The standard deviation is: 3.361811289755716

You have now mail in /var/spool/mail/acadgild